
Title: PM Accelerator TechAssessment Data Scientist

Author: Alex Ordonez

Date: March 21, 2025

Mission: By making industry-leading tools and education available to individuals from all backgrounds, we level the playing field for future PM leaders. This is the PM Accelerator motto, as we grant aspiring and experienced PMs what they need most – Access. We introduce you to industry leaders, surround you with the right PM ecosystem, and discover the new world of AI product management skills.

```
#Importing the necessary Libraries
import pandas as pd
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from pandas.api.types import is_numeric_dtype
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import shap
from sklearn.linear_model import LinearRegression
```

Analyze the "Global Weather Repository.csv" dataset to forecast future weather trends and showcase data science skills through both basic and advanced techniques.

```
#Reading and inspectioning data
df=pd.read_csv('./GlobalWeatherRepository.csv')
df.head()
```

	country	location_name	latitude	longitude	timezone
0	Afghanistan	Kabul	34.52	69.18	Asia/Kabul
1	Albania	Tirana	41.33	19.82	Europe/Tirane
2	Algeria	Algiers	36.76	3.05	Africa/Algiers
3	Andorra	Andorra La Vella	42.50	1.52	Europe/Andorra
4	Angola	Luanda	-8.84	13.23	Africa/Luanda

	last_updated_epoch	last_updated	temperature_celsius	\
0	1715849100	2024-05-16 13:15	26.6	
1	1715849100	2024-05-16 10:45	19.0	
2	1715849100	2024-05-16 09:45	23.0	
3	1715849100	2024-05-16 10:45	6.3	
4	1715849100	2024-05-16 09:45	26.0	

	temperature_fahrenheit	condition_text	...	air_quality_PM2.5	\
0	79.8	Partly Cloudy	...	8.4	
1	66.2	Partly cloudy	...	1.1	
2	73.4	Sunny	...	10.4	
3	43.3	Light drizzle	...	0.7	
4	78.8	Partly cloudy	...	183.4	

	air_quality_PM10	air_quality_us-epa-index	air_quality_gb-defra-index	\
0	26.6		1	
1				
1	2.0		1	
1				
2	18.4		1	
1				
3	0.9		1	
1				
4	262.3		5	
10				

	sunrise	sunset	moonrise	moonset	moon_phase
moon_illumination					
0	04:50 AM	06:50 PM	12:12 PM	01:11 AM	Waxing Gibbous
55					
1	05:21 AM	07:54 PM	12:58 PM	02:14 AM	Waxing Gibbous
55					
2	05:40 AM	07:50 PM	01:15 PM	02:14 AM	Waxing Gibbous
55					
3	06:31 AM	09:11 PM	02:12 PM	03:31 AM	Waxing Gibbous
55					
4	06:12 AM	05:55 PM	01:17 PM	12:38 AM	Waxing Gibbous
55					

[5 rows x 41 columns]

The first step that I took was to explore the data, verify types, null and nan values

```
#count of records
len(df)
```

```
59633
```

```
#Verifying columns
```

```
df.columns
```

```
Index(['country', 'location_name', 'latitude', 'longitude',  
      'timezone',  
      'last_updated_epoch', 'last_updated', 'temperature_celsius',  
      'temperature_fahrenheit', 'condition_text', 'wind_mph',  
      'wind_kph',  
      'wind_degree', 'wind_direction', 'pressure_mb', 'pressure_in',  
      'precip_mm', 'precip_in', 'humidity', 'cloud',  
      'feels_like_celsius',  
      'feels_like_fahrenheit', 'visibility_km', 'visibility_miles',  
      'uv_index', 'gust_mph', 'gust_kph',  
      'air_quality_Carbon_Monoxide',  
      'air_quality_Ozone', 'air_quality_Nitrogen_dioxide',  
      'air_quality_Sulphur_dioxide', 'air_quality_PM2.5',  
      'air_quality_PM10',  
      'air_quality_us-epa-index', 'air_quality_gb-defra-index',  
      'sunrise',  
      'sunset', 'moonrise', 'moonset', 'moon_phase',  
      'moon_illumination'],  
      dtype='object')
```

```
#Verifying types
```

```
df.dtypes
```

country	object
location_name	object
latitude	float64
longitude	float64
timezone	object
last_updated_epoch	int64
last_updated	object
temperature_celsius	float64
temperature_fahrenheit	float64
condition_text	object
wind_mph	float64
wind_kph	float64
wind_degree	int64
wind_direction	object
pressure_mb	float64
pressure_in	float64
precip_mm	float64
precip_in	float64
humidity	int64
cloud	int64
feels_like_celsius	float64
feels_like_fahrenheit	float64
visibility_km	float64
visibility_miles	float64

uv_index	float64
gust_mph	float64
gust_kph	float64
air_quality_Carbon_Monoxide	float64
air_quality_Ozone	float64
air_quality_Nitrogen_dioxide	float64
air_quality_Sulphur_dioxide	float64
air_quality_PM2.5	float64
air_quality_PM10	float64
air_quality_us-epa-index	int64
air_quality_gb-defra-index	int64
sunrise	object
sunset	object
moonrise	object
moonset	object
moon_phase	object
moon_illumination	int64
dtype:	object

#Localizing Null & Nan values

```
print("Null Values")
print(df.isnull().sum())
print("")
print("Nan Values")
print(df.isna().sum())
```

Null Values

country	0
location_name	0
latitude	0
longitude	0
timezone	0
last_updated_epoch	0
last_updated	0
temperature_celsius	0
temperature_fahrenheit	0
condition_text	0
wind_mph	0
wind_kph	0
wind_degree	0
wind_direction	0
pressure_mb	0
pressure_in	0
precip_mm	0
precip_in	0
humidity	0
cloud	0
feels_like_celsius	0
feels_like_fahrenheit	0
visibility_km	0

visibility_miles	0
uv_index	0
gust_mph	0
gust_kph	0
air_quality_Carbon_Monoxide	0
air_quality_Ozone	0
air_quality_Nitrogen_dioxide	0
air_quality_Sulphur_dioxide	0
air_quality_PM2.5	0
air_quality_PM10	0
air_quality_us-epa-index	0
air_quality_gb-defra-index	0
sunrise	0
sunset	0
moonrise	0
moonset	0
moon_phase	0
moon_illumination	0
dtype: int64	

Nan Values

country	0
location_name	0
latitude	0
longitude	0
timezone	0
last_updated_epoch	0
last_updated	0
temperature_celsius	0
temperature_fahrenheit	0
condition_text	0
wind_mph	0
wind_kph	0
wind_degree	0
wind_direction	0
pressure_mb	0
pressure_in	0
precip_mm	0
precip_in	0
humidity	0
cloud	0
feels_like_celsius	0
feels_like_fahrenheit	0
visibility_km	0
visibility_miles	0
uv_index	0
gust_mph	0
gust_kph	0
air_quality_Carbon_Monoxide	0

```

air_quality_Ozone          0
air_quality_Nitrogen_dioxide 0
air_quality_Sulphur_dioxide 0
air_quality_PM2.5          0
air_quality_PM10           0
air_quality_us-epa-index    0
air_quality_gb-defra-index  0
sunrise                    0
sunset                     0
moonrise                   0
moonset                    0
moon_phase                 0
moon_illumination          0
dtype: int64

```

#Observing Object string values

```
df[['sunrise', 'sunset', 'moonrise', 'moonset', 'moon_phase']]
```

	sunrise	sunset	moonrise	moonset	moon_phase
0	04:50 AM	06:50 PM	12:12 PM	01:11 AM	Waxing Gibbous
1	05:21 AM	07:54 PM	12:58 PM	02:14 AM	Waxing Gibbous
2	05:40 AM	07:50 PM	01:15 PM	02:14 AM	Waxing Gibbous
3	06:31 AM	09:11 PM	02:12 PM	03:31 AM	Waxing Gibbous
4	06:12 AM	05:55 PM	01:17 PM	12:38 AM	Waxing Gibbous
...
59628	06:32 AM	06:38 PM	10:58 PM	09:52 AM	Waning Gibbous
59629	06:02 AM	06:07 PM	10:24 PM	08:41 AM	Waning Gibbous
59630	06:08 AM	06:14 PM	10:27 PM	09:05 AM	Waning Gibbous
59631	06:11 AM	06:18 PM	09:34 PM	10:06 AM	Waning Gibbous
59632	06:00 AM	06:07 PM	09:17 PM	09:59 AM	Waning Gibbous

```
[59633 rows x 5 columns]
```

```
df=df.drop('last_updated_epoch', axis=1)
```

Calculating Measures for input variables is an important step and it would give me an overall idea of how is the data distributed

Also here I'm automating the outliers detection using IQR and Zscore, I used $(1.5 * IQR)$ to put the extreme fences and 3 std Zscore for detecting outliers.

#Calculating Important Measures for each numeric column

Getting IQR and Zscore to provide automatic detection of outliers

```
stats_df=pd.DataFrame({'Measure':
['Mean', 'Median', 'Std', 'Min', 'Max', 'Range', 'IQR', 'Outliers-IQR',
'Outliers-ZScore']})
```

```
for i in df.columns:
```

```
    #only numeric columns
```

```
    if is_numeric_dtype(df[i]):
```

```

Q3= df[i].quantile(0.75)
Q1= df[i].quantile(0.25)
IQR= Q3 - Q1

#Calculating ZScore
df['Zscore']=np.abs(stats.zscore(df[i]))

#Calculating Upper Fence and Lower Fence for IQR Outliers and
getting outliers
iqr_out=df[ (df[i] < Q1 - (1.5 * IQR)) | (df[i] > Q3 + (1.5 *
IQR))]

#Getting Outliers ZScore based on 3 std away from the mean
zcore_out=df[df['Zscore']>3]

# Calculating Measures
stats_df[i]=[df[i].mean(),
             df[i].median(),
             df[i].std(),
             df[i].min(),
             df[i].max(),
             df[i].max() - df[i].min(),
             IQR,
             len(iqr_out[i]),
             len(zcore_out[i])
            ]
print("Outliers IQR based:" + i )
print (iqr_out[i])
print("")
print("Outliers ZScore based:" + i )
print (zcore_out[i])
print("")

```

```

Outliers IQR based:latitude
Series([], Name: latitude, dtype: float64)

```

```

Outliers ZScore based:latitude
Series([], Name: latitude, dtype: float64)

```

```

Outliers IQR based:longitude
8      149.2200
58     178.4200
85     139.6900

```

```
89      169.5300
107     171.3800
...
59597   159.9500
59614   -175.2000
59619   179.2167
59624   -123.0439
59627   168.3167
Name: longitude, Length: 4588, dtype: float64
```

```
Outliers ZScore based:longitude
Series([], Name: longitude, dtype: float64)
```

```
Outliers IQR based:temperature_celsius
784      -1.0
3010     45.6
3205     45.7
4373     45.9
4568     46.6
...
59147    -1.7
59191    -2.9
59246    -0.1
59275    -5.7
59441    -2.0
Name: temperature_celsius, Length: 1389, dtype: float64
```

```
Outliers ZScore based:temperature_celsius
29235    -8.4
29820    -7.9
30600   -10.4
32884   -11.0
33079   -12.1
...
57133    -7.6
57328    -7.8
57577    -7.6
58495    -7.5
58661    -8.0
Name: temperature_celsius, Length: 306, dtype: float64
```

```
Outliers IQR based:temperature_fahrenheit
784      30.2
3205     114.3
4373     114.6
4568     115.9
5461      28.6
...
59063     30.6
59147     28.9
```



```
59191      26.8
59275      21.7
59441      28.4
Name: temperature_fahrenheit, Length: 1369, dtype: float64
```

Outliers ZScore based:temperature_fahrenheit

```
29235      16.9
29820      17.8
30600      13.2
32884      12.2
33079      10.3
```

```
...
57133      18.3
57328      18.0
57577      18.3
58495      18.5
58661      17.5
```

```
Name: temperature_fahrenheit, Length: 306, dtype: float64
```

Outliers IQR based:wind_mph

```
123        23.0
153        25.5
239        25.5
246        23.0
259        24.2
```

```
...
59504      22.6
59545      22.8
59550      26.2
59558      24.2
59578      25.1
```

```
Name: wind_mph, Length: 890, dtype: float64
```

Outliers ZScore based:wind_mph

```
834        106.9
1193       160.8
1827        40.5
3829        40.5
4997        36.7
7248       169.1
7317        36.7
7601      1841.2
8464        43.4
8659        37.8
8675       128.0
9129        50.3
12610       42.9
12897       40.5
13924       36.7
17747       36.7
```

19502	36.7
19892	48.5
20087	37.8
22816	47.2
33820	41.2
35380	47.6
39912	36.2
42278	38.5
43180	47.2
43375	39.4
43570	40.5
47658	42.3
48797	43.6
50161	39.1
50551	39.6
51136	37.1
51526	56.6
52696	38.7
52891	38.5
53475	40.0
53670	38.5
53864	36.9
54839	37.4
56593	41.8

Name: wind_mph, dtype: float64

Outliers IQR based:wind_kph

123	37.1
153	41.0
239	41.0
246	37.1
259	38.9

	...
59504	36.4
59545	36.7
59550	42.1
59558	38.9
59578	40.3

Name: wind_kph, Length: 890, dtype: float64

Outliers ZScore based:wind_kph

834	172.1
1193	258.8
1827	65.2
3829	65.2
4997	59.0
7248	272.2
7317	59.0
7601	2963.2

8464	69.8
8659	60.8
8675	205.9
9129	81.0
12610	69.1
12897	65.2
13924	59.0
17747	59.0
19502	59.0
19892	78.1
20087	60.8
22816	76.0
33820	66.2
35380	76.7
39912	58.3
42278	61.9
43180	76.0
43375	63.4
43570	65.2
47658	68.0
48797	70.2
50161	63.0
50551	63.7
51136	59.8
51526	91.1
52696	62.3
52891	61.9
53280	58.0
53475	64.4
53670	61.9
53864	59.4
54839	60.1
56593	67.3

Name: wind_kph, dtype: float64

Outliers IQR based:wind_degree

Series([], Name: wind_degree, dtype: int64)

Outliers ZScore based:wind_degree

Series([], Name: wind_degree, dtype: int64)

Outliers IQR based:pressure_mb

119	1033.0
161	1031.0
1894	1031.0
1952	1031.0
2154	997.0

	...
59580	1033.0
59587	1033.0

```
59591    1035.0
59595    1032.0
59596    1033.0
Name: pressure_mb, Length: 2302, dtype: float64
```

Outliers ZScore based:pressure_mb

```
30172    971.0
38510    1080.0
45744    972.0
48992    973.0
49582    964.0
49769    3006.0
50941    947.0
52114    3000.0
54449    964.0
54644    969.0
54839    962.0
55034    964.0
```

Name: pressure_mb, dtype: float64

Outliers IQR based:pressure_in

```
95      30.39
119     30.50
161     30.45
314     30.39
1764    29.46
```

```
...
59580    30.50
59587    30.50
59591    30.56
59595    30.47
59596    30.50
```

Name: pressure_in, Length: 3365, dtype: float64

Outliers ZScore based:pressure_in

```
30172    28.67
38510    31.89
45744    28.70
48992    28.73
49582    28.47
49769    88.77
50941    27.96
52114    88.59
54449    28.47
54644    28.61
54839    28.41
55034    28.47
```

Name: pressure_in, dtype: float64

Outliers IQR based:precip_mm

```
1      0.10
3      0.30
7      0.13
16     0.25
19     0.16
...
59610  0.40
59612  0.08
59619  0.12
59620  0.27
59628  0.84
Name: precip_mm, Length: 11167, dtype: float64
```

Outliers ZScore based:precip_mm

```
176    2.09
181    2.00
219    3.01
273    2.05
317    2.42
```

```
...
59352  1.97
59378  2.35
59417  4.01
59582  3.20
59594  8.52
```

Name: precip_mm, Length: 762, dtype: float64

Outliers IQR based:precip_in

```
3      0.01
7      0.01
16     0.01
19     0.01
40     0.01
```

```
...
59597  0.02
59599  0.02
59610  0.02
59620  0.01
59628  0.03
```

Name: precip_in, Length: 9371, dtype: float64

Outliers ZScore based:precip_in

```
176    0.08
181    0.08
219    0.12
273    0.08
317    0.10
```

```
...
59352  0.08
59378  0.09
```

```

59417    0.16
59582    0.13
59594    0.34
Name: precip_in, Length: 807, dtype: float64

Outliers IQR based:humidity
Series([], Name: humidity, dtype: int64)

Outliers ZScore based:humidity
Series([], Name: humidity, dtype: int64)

Outliers IQR based:cloud
Series([], Name: cloud, dtype: int64)

Outliers ZScore based:cloud
Series([], Name: cloud, dtype: int64)

Outliers IQR based:feels_like_celsius
6816     -4.9
9123     -4.3
28066    -7.1
28612    -5.3
28807    -7.4
...
59280    -4.6
59342    -5.2
59426    -4.4
59441    -5.4
59581    -4.1
Name: feels_like_celsius, Length: 1311, dtype: float64

Outliers ZScore based:feels_like_celsius
29040    -13.6
29820    -12.6
30405    -14.6
30600    -14.8
32884    -19.2
...
57188    -13.9
57328    -15.5
57577    -14.6
57772    -14.2
58661    -12.7
Name: feels_like_celsius, Length: 314, dtype: float64

Outliers IQR based:feels_like_fahrenheit
6816     23.1
9123     24.2
14117    25.1
28066    19.3

```

```
28612    22.4
...
59280    23.7
59342    22.6
59426    24.1
59441    22.2
59581    24.6
Name: feels_like_fahrenheit, Length: 1325, dtype: float64
```

Outliers ZScore based:feels_like_fahrenheit

```
29040    7.5
29820    9.4
30405    5.8
30600    5.3
32884   -2.5
```

```
...
57188    7.1
57328    4.1
57577    5.7
57772    6.5
58661    9.1
```

Name: feels_like_fahrenheit, Length: 315, dtype: float64

Outliers IQR based:visibility_km

```
3         2.0
23        0.0
32       24.0
35        7.0
40        7.0
```

```
...
59597     9.0
59599     9.0
59610     9.0
59624    16.0
59628     9.0
```

Name: visibility_km, Length: 11632, dtype: float64

Outliers ZScore based:visibility_km

```
23        0.0
32       24.0
107       24.0
111       24.0
132       23.0
```

```
...
59375    21.0
59461     0.0
59470    24.0
59549    19.0
59570    21.0
```

Name: visibility_km, Length: 1620, dtype: float64

```
Outliers IQR based:visibility_miles
3      1.0
23     0.0
32    14.0
35     4.0
40     4.0
...
59597   5.0
59599   5.0
59610   5.0
59624   9.0
59628   5.0
Name: visibility_miles, Length: 11541, dtype: float64
```

```
Outliers ZScore based:visibility_miles
23     0.0
32    14.0
107    14.0
111    14.0
132    14.0
...
59375  13.0
59461   0.0
59470  14.0
59549  11.0
59570  13.0
Name: visibility_miles, Length: 1585, dtype: float64
```

```
Outliers IQR based:uv_index
Series([], Name: uv_index, dtype: float64)
```

```
Outliers ZScore based:uv_index
37958  15.3
39689  15.2
40281  15.0
40567  15.0
40568  15.2
...
58746  15.8
58941  15.2
59220  15.2
59415  14.9
59610  15.0
Name: uv_index, Length: 189, dtype: float64
```

```
Outliers IQR based:gust_mph
85     32.5
153    30.0
176    35.3
```



```
207      38.5
239      30.0
...
59546    30.0
59550    30.1
59561    30.5
59578    36.9
59590    31.3
Name: gust_mph, Length: 1365, dtype: float64
```

Outliers ZScore based:gust_mph

```
512      48.0
834     111.4
1193     165.3
1465      43.3
1827      69.8
...
56398     47.3
56593     53.3
58569     51.7
58976     48.2
59058     43.7
Name: gust_mph, Length: 128, dtype: float64
```

Outliers IQR based:gust_kph

```
85       52.2
153      48.2
176      56.8
207      62.0
239      48.2
...
59546     48.3
59550     48.5
59561     49.1
59578     59.4
59590     50.4
Name: gust_kph, Length: 1347, dtype: float64
```

Outliers ZScore based:gust_kph

```
512      77.3
834     179.3
1193     266.0
1465      69.6
1827     112.3
...
56398     76.1
56593     85.8
58569     83.2
58976     77.6
59058     70.3
```

Name: gust_kph, Length: 128, dtype: float64

Outliers IQR based:air_quality_Carbon_Monoxide

4	2964.00
30	1295.10
35	2723.70
36	1335.10
40	1161.60

	...
59575	1036.00
59600	1052.65
59603	1585.45
59611	1359.75
59628	2099.75

Name: air_quality_Carbon_Monoxide, Length: 5220, dtype: float64

Outliers ZScore based:air_quality_Carbon_Monoxide

78	3471.400
173	3898.600
230	9719.900
245	3845.200
273	19653.301

	...
59128	5078.250
59279	4253.150
59321	4774.850
59474	3494.650
59516	6184.550

Name: air_quality_Carbon_Monoxide, Length: 761, dtype: float64

Outliers IQR based:air_quality_Ozone

13	188.8
78	303.3
80	161.7
88	160.2
130	173.1

	...
59213	206.0
59355	151.0
59408	167.0
59559	155.0
59603	182.0

Name: air_quality_Ozone, Length: 1270, dtype: float64

Outliers ZScore based:air_quality_Ozone

13	188.8
78	303.3
130	173.1
191	197.4
275	176.0

```
...
58001    177.0
58433    180.0
58823    192.0
59213    206.0
59603    182.0
Name: air_quality_Ozone, Length: 596, dtype: float64
```

```
Outliers IQR based:air_quality_Nitrogen_dioxide
```

```
2         65.100
4         72.700
35        41.800
36       101.500
50        72.700
```

```
...
59602     73.815
59607     40.515
59618     51.615
59622     70.115
59623     83.250
```

```
Name: air_quality_Nitrogen_dioxide, Length: 7474, dtype: float64
```

```
Outliers ZScore based:air_quality_Nitrogen_dioxide
```

```
36       101.500
79       145.300
103      133.000
230      181.000
231      128.900
```

```
...
59516    116.735
59517    117.475
59541    114.515
59579     95.645
59600     96.385
```

```
Name: air_quality_Nitrogen_dioxide, Length: 1501, dtype: float64
```

```
Outliers IQR based:air_quality_Sulphur_dioxide
```

```
4         31.500
35        24.800
36       223.200
50        52.000
78        40.100
```

```
...
59603     76.960
59608     23.125
59611     25.530
59622     38.665
59629     29.415
```

```
Name: air_quality_Sulphur_dioxide, Length: 8433, dtype: float64
```

Outliers ZScore based:air_quality_Sulphur_dioxide

36	223.200
231	165.900
1007	177.400
1440	200.300
1787	213.600

	...
58042	165.945
58153	178.155
58169	187.960
58320	181.485
59218	159.470

Name: air_quality_Sulphur_dioxide, Length: 354, dtype: float64

Outliers IQR based:air_quality_PM2.5

4	183.400
35	211.100
36	84.900
68	132.000
78	196.100

	...
59548	81.030
59589	253.265
59590	66.785
59600	78.625
59601	107.855

Name: air_quality_PM2.5, Length: 4665, dtype: float64

Outliers ZScore based:air_quality_PM2.5

4	183.400
35	211.100
78	196.100
230	714.100
231	228.200

	...
58571	186.850
59004	234.395
59474	189.995
59516	159.840
59589	253.265

Name: air_quality_PM2.5, Length: 758, dtype: float64

Outliers IQR based:air_quality_PM10

4	262.300
12	114.300
35	268.600
36	107.800
68	178.100

	...
59589	2729.120

```
59590      289.525
59601      627.150
59604      105.080
59630      173.530
Name: air_quality_PM10, Length: 6001, dtype: float64
```

Outliers ZScore based:air_quality_PM10

```
230      873.400
273      621.500
468     1002.200
617      682.100
811      544.400
```

```
...
59211     628.815
59351     646.020
59406     673.585
59589    2729.120
59601     627.150
```

Name: air_quality_PM10, Length: 609, dtype: float64

Outliers IQR based:air_quality_us-epa-index

```
4         5
35        5
36        4
68        4
78        5
```

```
..
59548     4
59589     6
59590     4
59600     4
59601     4
```

Name: air_quality_us-epa-index, Length: 4726, dtype: int64

Outliers ZScore based:air_quality_us-epa-index

```
4         5
35        5
78        5
230       6
231       5
```

```
..
58571     5
59004     5
59474     5
59516     5
59589     6
```

Name: air_quality_us-epa-index, Length: 851, dtype: int64

Outliers IQR based:air_quality_gb-defra-index

```
4         10
```

```

35      10
36      10
68      10
78      10
..
59590    9
59600    10
59601    10
59611     8
59622     8
Name: air_quality_gb-defra-index, Length: 6563, dtype: int64

Outliers ZScore based:air_quality_gb-defra-index
Series([], Name: air_quality_gb-defra-index, dtype: int64)

Outliers IQR based:moon_illumination
Series([], Name: moon_illumination, dtype: int64)

Outliers ZScore based:moon_illumination
Series([], Name: moon_illumination, dtype: int64)

stats_df.style
<pandas.io.formats.style.Styler at 0x220a0236760>

```

Here I wanted to add certain columns dividing the date getting the month year, weekday and hour that I believe that if the moment in which an observation is taken is really important depending of the month or hour.

Also I am scaling all the numeric values to be between 0-1

```

#Scaling and Adding some columns
scaler=MinMaxScaler()

#Transform 'last_updated' to datetime
df['last_updated']=pd.to_datetime(df['last_updated'])

# Getting date time columns
df['year']=df['last_updated'].dt.year
df['month']=df['last_updated'].dt.month
df['day']=df['last_updated'].dt.day
df['weekday']=df['last_updated'].dt.day_name()
df['hour']=df['last_updated'].dt.hour

# Transforming Month and Hour to catch the cyclic relationship between hours
df['Hour_sin'] = np.sin(2 * np.pi * df['hour'] / 24)
df['Hour_cos'] = np.cos(2 * np.pi * df['hour'] / 24)
df['Month_sin'] = np.sin(2 * np.pi * df['month'] / 12)
df['Month_cos'] = np.cos(2 * np.pi * df['month'] / 12)

```

```

#Columns to scale
columns = stats_df.columns[1:].tolist()
columns.append('Hour_sin')
columns.append('Hour_cos')
columns.append('Month_sin')
columns.append('Month_cos')

#Scaling values
df_numeric_normalized =
pd.DataFrame(scaler.fit_transform(df[columns]), columns=columns)

df_numeric_normalized

```

	latitude	longitude	temperature_celsius
temperature_fahrenheit \			
0	0.719014	0.689521	0.695007
0.694153			
1	0.783594	0.550251	0.592443
0.592204			
2	0.740256	0.502934	0.646424
0.646177			
3	0.794689	0.498617	0.421053
0.420540			
4	0.307824	0.531657	0.686910
0.686657			
...
.			..
59628	0.491228	0.305523	0.676113
0.676162			
59629	0.591117	0.792986	0.636977
0.636432			
59630	0.537266	0.619058	0.632928
0.631934			
59631	0.245456	0.574130	0.636977
0.637181			
59632	0.222686	0.581922	0.639676
0.640180			

	wind_mph	wind_kph	wind_degree	pressure_mb	pressure_in
precip_mm \					
0	0.003317	0.003277	0.938719	0.031569	0.031738
0.000000					
1	0.002556	0.002568	0.888579	0.031569	0.031574
0.002367					
2	0.003915	0.003886	0.777159	0.031083	0.031080
0.000000					
3	0.002828	0.002804	0.596100	0.029140	0.029436

0.007102					
4	0.003208	0.003176	0.415042	0.031083	0.031080
0.000000					
...
...					
59628	0.000000	0.000000	0.136490	0.032054	0.032067
0.019886					
59629	0.002828	0.002804	0.977716	0.036911	0.037000
0.000000					
59630	0.000272	0.000237	0.529248	0.033511	0.033547
0.000000					
59631	0.004296	0.004257	0.225627	0.033997	0.034205
0.000237					
59632	0.001468	0.001453	0.178273	0.034483	0.034534
0.000000					

	air_quality_Sulphur_dioxide	air_quality_PM2.5
air_quality_PM10 \		
0	0.950464	0.005090
0.237748		
1	0.950455	0.000567
0.234629		
2	0.951719	0.006329
0.236708		
3	0.950464	0.000319
0.234489		
4	0.953440	0.113522
0.267639		
...
...		
59628	0.950727	0.010890
0.237355		
59629	0.953241	0.022123
0.239771		
59630	0.950498	0.017767
0.256381		
59631	0.950516	0.003324
0.235220		
59632	0.950498	0.005846
0.235759		

	air_quality_us-epa-index	air_quality_gb-defra-index \
0	0.0	0.000000
1	0.0	0.000000
2	0.0	0.000000
3	0.0	0.000000
4	0.8	1.000000
...
59628	0.2	0.111111

59629	0.2	0.222222
59630	0.2	0.222222
59631	0.0	0.000000
59632	0.0	0.000000

	moon_illumination	Hour_sin	Hour_cos	Month_sin	Month_cos
0	0.55	0.370590	0.017037	0.75	0.066987
1	0.55	0.750000	0.066987	0.75	0.066987
2	0.55	0.853553	0.146447	0.75	0.066987
3	0.55	0.750000	0.066987	0.75	0.066987
4	0.55	0.853553	0.146447	0.75	0.066987
...
59628	0.80	0.982963	0.629410	1.00	0.500000
59629	0.83	0.066987	0.250000	1.00	0.500000
59630	0.82	0.500000	0.000000	1.00	0.500000
59631	0.82	0.629410	0.017037	1.00	0.500000
59632	0.82	0.629410	0.017037	1.00	0.500000

[59633 rows x 33 columns]

```
#Combining Categorical and Numerical columns
df_scaled = df.select_dtypes(exclude=['number'])
df_scaled[colums]= df_numeric_normalized
df_scaled
```

	country	location_name	timezone	
last_updated \				
0	Afghanistan	Kabul	Asia/Kabul	2024-05-16 13:15:00
1	Albania	Tirana	Europe/Tirane	2024-05-16 10:45:00
2	Algeria	Algiers	Africa/Algiers	2024-05-16 09:45:00
3	Andorra	Andorra La Vella	Europe/Andorra	2024-05-16 10:45:00
4	Angola	Luanda	Africa/Luanda	2024-05-16 09:45:00
...
...				
59628	Venezuela	Caracas	America/Caracas	2025-03-19 05:30:00
59629	Vietnam	Hanoi	Asia/Bangkok	2025-03-19 16:30:00
59630	Yemen	Sanaa	Asia/Aden	2025-03-19 12:45:00
59631	Zambia	Lusaka	Africa/Lusaka	2025-03-19 11:45:00
59632	Zimbabwe	Harare	Africa/Harare	2025-03-19 11:30:00

	condition_text	wind_direction	sunrise	sunset	moonrise
\					
0	Partly Cloudy	NNW	04:50 AM	06:50 PM	12:12 PM
1	Partly cloudy	NW	05:21 AM	07:54 PM	12:58 PM
2	Sunny	W	05:40 AM	07:50 PM	01:15 PM
3	Light drizzle	SW	06:31 AM	09:11 PM	02:12 PM
4	Partly cloudy	SSE	06:12 AM	05:55 PM	01:17 PM
...
59628	Clear	NE	06:32 AM	06:38 PM	10:58 PM
59629	Sunny	N	06:02 AM	06:07 PM	10:24 PM
59630	Sunny	SSW	06:08 AM	06:14 PM	10:27 PM
59631	Patchy rain nearby	E	06:11 AM	06:18 PM	09:34 PM
59632	Sunny	ENE	06:00 AM	06:07 PM	09:17 PM
	moonset	...	air_quality_Sulphur_dioxide	air_quality_PM2.5	\
0	01:11 AM	...	0.950464	0.005090	
1	02:14 AM	...	0.950455	0.000567	
2	02:14 AM	...	0.951719	0.006329	
3	03:31 AM	...	0.950464	0.000319	
4	12:38 AM	...	0.953440	0.113522	
...	
59628	09:52 AM	...	0.950727	0.010890	
59629	08:41 AM	...	0.953241	0.022123	
59630	09:05 AM	...	0.950498	0.017767	
59631	10:06 AM	...	0.950516	0.003324	
59632	09:59 AM	...	0.950498	0.005846	
	air_quality_PM10	air_quality_us-epa-index	air_quality_gb-defra-index	\	
0	0.237748	0.0			
0.000000					
1	0.234629	0.0			
0.000000					
2	0.236708	0.0			
0.000000					
3	0.234489	0.0			
0.000000					
4	0.267639	0.8			
1.000000					

```

...
...
59628      0.237355      0.2
0.111111
59629      0.239771      0.2
0.222222
59630      0.256381      0.2
0.222222
59631      0.235220      0.0
0.000000
59632      0.235759      0.0
0.000000

      moon_illumination  Hour_sin  Hour_cos  Month_sin  Month_cos
0      0.55  0.370590  0.017037      0.75  0.066987
1      0.55  0.750000  0.066987      0.75  0.066987
2      0.55  0.853553  0.146447      0.75  0.066987
3      0.55  0.750000  0.066987      0.75  0.066987
4      0.55  0.853553  0.146447      0.75  0.066987
...
59628      0.80  0.982963  0.629410      1.00  0.500000
59629      0.83  0.066987  0.250000      1.00  0.500000
59630      0.82  0.500000  0.000000      1.00  0.500000
59631      0.82  0.629410  0.017037      1.00  0.500000
59632      0.82  0.629410  0.017037      1.00  0.500000

[59633 rows x 45 columns]

```

When I was exploring the data I notice that the columns of sunset, sunrise, moonrise, moonset were in a string format not 0 - 23 so here I am paring those columns to be 0-23

#Parsing Hour of the day for sunrise

```

df_scaled['split1']=df_scaled['sunrise'].str[:2]
df_scaled['split2']=df_scaled['sunrise'].str[-2:]
df_scaled['sunrise']=df_scaled.apply(lambda x: np.nan if not
x['split1'].isnumeric() else

(int(x['split1']) if x['split2']=='AM' and int(x['split1'])!=12 else

(0 if x['split2']=='AM' and int(x['split1'])==12 else

(int(x['split1'])+12 if int(x['split1'])!=12

else 12))),axis=1)
df_scaled=df_scaled.drop('split1', axis=1)
df_scaled=df_scaled.drop('split2', axis=1)

```

```

#Parsing Hour of the day for sunset

df_scaled['split1']=df_scaled['sunset'].str[:2]
df_scaled['split2']=df_scaled['sunset'].str[-2:]
df_scaled['sunset']=df_scaled.apply(lambda x: np.nan if not
x['split1'].isnumeric() else

(int(x['split1']) if x['split2']=='AM' and int(x['split1'])!=12 else

(0 if x['split2']=='AM' and int(x['split1'])==12 else

(int(x['split1'])+12 if int(x['split1'])!=12

else 12))),axis=1)
df_scaled=df_scaled.drop('split1', axis=1)
df_scaled=df_scaled.drop('split2', axis=1)

#Parsing Hour of the day for moonrise

df_scaled['split1']=df_scaled['moonrise'].str[:2]
df_scaled['split2']=df_scaled['moonrise'].str[-2:]
df_scaled['moonrise']=df_scaled.apply(lambda x: np.nan if not
x['split1'].isnumeric() else

(int(x['split1']) if x['split2']=='AM' and int(x['split1'])!=12 else

(0 if x['split2']=='AM' and int(x['split1'])==12 else

(int(x['split1'])+12 if int(x['split1'])!=12

else 12))),axis=1)
df_scaled=df_scaled.drop('split1', axis=1)
df_scaled=df_scaled.drop('split2', axis=1)

#Parsing Hour of the day for moonset

df_scaled['split1']=df_scaled['moonset'].str[:2]
df_scaled['split2']=df_scaled['moonset'].str[-2:]
df_scaled['moonset']=df_scaled.apply(lambda x: np.nan if not
x['split1'].isnumeric() else

(int(x['split1']) if x['split2']=='AM' and int(x['split1'])!=12 else

(0 if x['split2']=='AM' and int(x['split1'])==12 else

(int(x['split1'])+12 if int(x['split1'])!=12

else 12))),axis=1)
df_scaled=df_scaled.drop('split1', axis=1)
df_scaled=df_scaled.drop('split2', axis=1)

```

df_scaled

	country	location_name	timezone	
last_updated \				
0	Afghanistan	Kabul	Asia/Kabul	2024-05-16
13:15:00				
1	Albania	Tirana	Europe/Tirane	2024-05-16
10:45:00				
2	Algeria	Algiers	Africa/Algiers	2024-05-16
09:45:00				
3	Andorra	Andorra La Vella	Europe/Andorra	2024-05-16
10:45:00				
4	Angola	Luanda	Africa/Luanda	2024-05-16
09:45:00				
...	
...				
59628	Venezuela	Caracas	America/Caracas	2025-03-19
05:30:00				
59629	Vietnam	Hanoi	Asia/Bangkok	2025-03-19
16:30:00				
59630	Yemen	Sanaa	Asia/Aden	2025-03-19
12:45:00				
59631	Zambia	Lusaka	Africa/Lusaka	2025-03-19
11:45:00				
59632	Zimbabwe	Harare	Africa/Harare	2025-03-19
11:30:00				

	condition_text	wind_direction	sunrise	sunset	moonrise
moonset \					
0	Partly Cloudy	NNW	4	18	12.0
1.0					
1	Partly cloudy	NW	5	19	12.0
2.0					
2	Sunny	W	5	19	13.0
2.0					
3	Light drizzle	SW	6	21	14.0
3.0					
4	Partly cloudy	SSE	6	17	13.0
0.0					
...
...					
59628	Clear	NE	6	18	22.0
9.0					
59629	Sunny	N	6	18	22.0
8.0					
59630	Sunny	SSW	6	18	22.0
9.0					
59631	Patchy rain nearby	E	6	18	21.0
10.0					
59632	Sunny	ENE	6	18	21.0

9.0

	...	air_quality_Sulphur_dioxide	air_quality_PM2.5
air_quality_PM10 \			
0	...	0.950464	0.005090
0.237748			
1	...	0.950455	0.000567
0.234629			
2	...	0.951719	0.006329
0.236708			
3	...	0.950464	0.000319
0.234489			
4	...	0.953440	0.113522
0.267639			
...
...			
59628	...	0.950727	0.010890
0.237355			
59629	...	0.953241	0.022123
0.239771			
59630	...	0.950498	0.017767
0.256381			
59631	...	0.950516	0.003324
0.235220			
59632	...	0.950498	0.005846
0.235759			

	air_quality_us-epa-index	air_quality_gb-defra-index \
0	0.0	0.000000
1	0.0	0.000000
2	0.0	0.000000
3	0.0	0.000000
4	0.8	1.000000
...
59628	0.2	0.111111
59629	0.2	0.222222
59630	0.2	0.222222
59631	0.0	0.000000
59632	0.0	0.000000

	moon_illumination	Hour_sin	Hour_cos	Month_sin	Month_cos
0	0.55	0.370590	0.017037	0.75	0.066987
1	0.55	0.750000	0.066987	0.75	0.066987
2	0.55	0.853553	0.146447	0.75	0.066987
3	0.55	0.750000	0.066987	0.75	0.066987
4	0.55	0.853553	0.146447	0.75	0.066987
...
59628	0.80	0.982963	0.629410	1.00	0.500000
59629	0.83	0.066987	0.250000	1.00	0.500000
59630	0.82	0.500000	0.000000	1.00	0.500000

59631	0.82	0.629410	0.017037	1.00	0.500000
59632	0.82	0.629410	0.017037	1.00	0.500000

[59633 rows x 45 columns]

Here I start plotting, correlation matrix to review correlations, of course there are some variables that I expect to be highly correlated like temperature_celsius and temperature_fahrenheit

#Building a correlation Matrix

```
correlation_matrix = df_scaled[columns].corr()
correlation_matrix
```

	latitude	longitude	temperature_celsius
latitude	1.000000	-0.020472	-0.360750
longitude	-0.020472	1.000000	0.092861
temperature_celsius	-0.360750	0.092861	1.000000
temperature_fahrenheit	-0.360738	0.092847	0.999997
wind_mph	0.020455	0.024700	0.065894
wind_kph	0.020485	0.024626	0.065820
wind_degree	0.166279	0.064298	-0.047451
pressure_mb	0.055717	-0.076264	-0.274208
pressure_in	0.055882	-0.075219	-0.274797
precip_mm	-0.053977	0.054391	0.020900
precip_in	-0.052638	0.054409	0.020956
humidity	-0.075062	-0.177452	-0.345254
cloud	-0.048952	0.014072	-0.174460
feels_like_celsius	-0.384608	0.093044	0.981122
feels_like_fahrenheit	-0.384603	0.093047	0.981117
visibility_km	-0.030959	0.129263	0.089081
visibility_miles	-0.032864	0.128075	0.093096
uv_index	-0.176908	-0.023290	0.531248

gust_mph	0.005723	0.020047	0.080642
gust_kph	0.005746	0.020032	0.080661
air_quality_Carbon_Monoxide	-0.036401	0.106860	-0.083877
air_quality_Ozone	0.105086	0.033032	0.288131
air_quality_Nitrogen_dioxide	0.243139	0.138328	-0.287191
air_quality_Sulphur_dioxide	0.056135	0.077484	-0.081419
air_quality_PM2.5	-0.008816	0.071967	-0.061818
air_quality_PM10	-0.001675	0.030109	0.039651
air_quality_us-epa-index	0.060135	0.126316	-0.057699
air_quality_gb-defra-index	0.061964	0.122535	-0.047209
moon_illumination	-0.000289	-0.001131	-0.014185
Hour_sin	-0.117487	-0.703979	-0.233213
Hour_cos	-0.291952	0.140608	0.076940
Month_sin	0.000908	-0.001179	-0.229576
Month_cos	-0.000320	0.003569	-0.354801
		temperature_fahrenheit	wind_mph
wind_kph \			
latitude		-0.360738	0.020455
0.020485			
longitude		0.092847	0.024700
0.024626			
temperature_celsius		0.999997	0.065894
0.065820			
temperature_fahrenheit		1.000000	0.065910
0.065836			
wind_mph		0.065910	1.000000
0.999992			
wind_kph		0.065836	0.999992
1.000000			
wind_degree		-0.047436	0.003051
0.003043			
pressure_mb		-0.274201	-0.055280 -
0.055250			
pressure_in		-0.274791	-0.055476 -
0.055445			

precip_mm	0.020910	0.001028
0.001002		
precip_in	0.020966	0.000430
0.000403		
humidity	-0.345254	-0.065447 -
0.065415		
cloud	-0.174451	0.009650
0.009661		
feels_like_celsius	0.981122	0.048397
0.048339		
feels_like_fahrenheit	0.981117	0.048416
0.048358		
visibility_km	0.089086	0.055764
0.055695		
visibility_miles	0.093099	0.057769
0.057707		
uv_index	0.531252	0.049038
0.048987		
gust_mph	0.080655	0.952600
0.952626		
gust_kph	0.080674	0.952597
0.952623		
air_quality_Carbon_Monoxide	-0.083877	-0.087321 -
0.087301		
air_quality_Ozone	0.288146	0.070747
0.070747		
air_quality_Nitrogen_dioxide	-0.287183	-0.108204 -
0.108187		
air_quality_Sulphur_dioxide	-0.081417	-0.043374 -
0.043360		
air_quality_PM2.5	-0.061810	-0.046413 -
0.046441		
air_quality_PM10	0.039658	0.033776
0.033713		
air_quality_us-epa-index	-0.057699	-0.065643 -
0.065681		
air_quality_gb-defra-index	-0.047210	-0.058086 -
0.058119		
moon_illumination	-0.014187	0.008497
0.008491		
Hour_sin	-0.233204	-0.063190 -
0.063147		
Hour_cos	0.076925	-0.029292 -
0.029262		
Month_sin	-0.229575	0.007970
0.007973		
Month_cos	-0.354802	-0.043862 -
0.043886		

	wind_degree	pressure_mb	pressure_in	\
latitude	0.166279	0.055717	0.055882	
longitude	0.064298	-0.076264	-0.075219	
temperature_celsius	-0.047451	-0.274208	-0.274797	
temperature_fahrenheit	-0.047436	-0.274201	-0.274791	
wind_mph	0.003051	-0.055280	-0.055476	
wind_kph	0.003043	-0.055250	-0.055445	
wind_degree	1.000000	-0.035822	-0.035533	
pressure_mb	-0.035822	1.000000	0.999873	
pressure_in	-0.035533	0.999873	1.000000	
precip_mm	0.012641	-0.062187	-0.061979	
precip_in	0.013013	-0.061343	-0.061138	
humidity	-0.037420	0.006319	0.006198	
cloud	0.010097	-0.024254	-0.024236	
feels_like_celsius	-0.053672	-0.273277	-0.273845	
feels_like_fahrenheit	-0.053649	-0.273279	-0.273847	
visibility_km	-0.077087	-0.014825	-0.015371	
visibility_miles	-0.077014	-0.013735	-0.014210	
uv_index	0.027623	-0.073106	-0.073561	
gust_mph	-0.001167	-0.082262	-0.082546	
gust_kph	-0.001144	-0.082263	-0.082546	
air_quality_Carbon_Monoxide	0.018423	0.026819	0.027279	
air_quality_Ozone	0.072918	-0.105671	-0.105882	
air_quality_Nitrogen_dioxide	0.067661	0.101824	0.102609	
air_quality_Sulphur_dioxide	0.040488	0.039040	0.039241	
air_quality_PM2.5	-0.009652	0.037686	0.037811	
air_quality_PM10	-0.020724	0.008870	0.008719	
air_quality_us-epa-index	-0.000689	0.054726	0.054910	
air_quality_gb-defra-index	-0.000876	0.048552	0.048720	
moon_illumination	-0.015165	0.004325	0.004248	
Hour_sin	-0.171327	0.110166	0.109265	
Hour_cos	-0.113386	-0.116270	-0.115996	
Month_sin	-0.038939	0.068977	0.068855	
Month_cos	-0.066503	0.124114	0.124354	
	precip_mm	...		
air_quality_Sulphur_dioxide	\			
latitude	-0.053977	...		
0.056135				
longitude	0.054391	...		
0.077484				
temperature_celsius	0.020900	...	-	
0.081419				
temperature_fahrenheit	0.020910	...	-	
0.081417				
wind_mph	0.001028	...	-	
0.043374				
wind_kph	0.001002	...	-	
0.043360				

wind_degree	0.012641	...	
0.040488			
pressure_mb	-0.062187	...	
0.039040			
pressure_in	-0.061979	...	
0.039241			
precip_mm	1.000000	...	-
0.020338			
precip_in	0.998237	...	-
0.018903			
humidity	0.183946	...	-
0.063358			
cloud	0.214786	...	-
0.070107			
feels_like_celsius	0.049077	...	-
0.079585			
feels_like_fahrenheit	0.049081	...	-
0.079607			
visibility_km	-0.046091	...	-
0.036296			
visibility_miles	-0.056418	...	-
0.039468			
uv_index	-0.067069	...	-
0.071125			
gust_mph	0.041330	...	-
0.059642			
gust_kph	0.041318	...	-
0.059666			
air_quality_Carbon_Monoxide	0.009413	...	
0.275804			
air_quality_Ozone	-0.086634	...	-
0.030622			
air_quality_Nitrogen_dioxide	-0.027682	...	
0.366879			
air_quality_Sulphur_dioxide	-0.020338	...	
1.000000			
air_quality_PM2.5	-0.047968	...	
0.252439			
air_quality_PM10	-0.042026	...	
0.114911			
air_quality_us-epa-index	-0.074251	...	
0.283694			
air_quality_gb-defra-index	-0.070285	...	
0.280532			
moon_illumination	0.005482	...	-
0.001131			
Hour_sin	-0.044366	...	-
0.065229			
Hour_cos	0.076860	...	-

0.014875	
Month_sin	-0.018893 ...
0.042896	
Month_cos	-0.028547 ...
0.090919	

	air_quality_PM2.5	air_quality_PM10 \
latitude	-0.008816	-0.001675
longitude	0.071967	0.030109
temperature_celsius	-0.061818	0.039651
temperature_fahrenheit	-0.061810	0.039658
wind_mph	-0.046413	0.033776
wind_kph	-0.046441	0.033713
wind_degree	-0.009652	-0.020724
pressure_mb	0.037686	0.008870
pressure_in	0.037811	0.008719
precip_mm	-0.047968	-0.042026
precip_in	-0.046252	-0.040685
humidity	-0.135112	-0.194318
cloud	-0.160064	-0.150691
feels_like_celsius	-0.068467	0.020729
feels_like_fahrenheit	-0.068490	0.020719
visibility_km	-0.123629	-0.059455
visibility_miles	-0.134258	-0.065278
uv_index	-0.066759	0.036126
gust_mph	-0.071034	0.020886
gust_kph	-0.071053	0.020879
air_quality_Carbon_Monoxide	0.632877	0.200310
air_quality_Ozone	0.009886	0.046597
air_quality_Nitrogen_dioxide	0.486304	0.188252
air_quality_Sulphur_dioxide	0.252439	0.114911
air_quality_PM2.5	1.000000	0.627611
air_quality_PM10	0.627611	1.000000
air_quality_us-epa-index	0.769686	0.513657
air_quality_gb-defra-index	0.720709	0.480097
moon_illumination	-0.002465	-0.001317
Hour_sin	-0.005076	0.009843
Hour_cos	-0.080553	-0.103605
Month_sin	0.104827	0.072800
Month_cos	0.160459	0.106816

	air_quality_us-epa-index \
latitude	0.060135
longitude	0.126316
temperature_celsius	-0.057699
temperature_fahrenheit	-0.057699
wind_mph	-0.065643
wind_kph	-0.065681
wind_degree	-0.000689

pressure_mb	0.054726
pressure_in	0.054910
precip_mm	-0.074251
precip_in	-0.071787
humidity	-0.196531
cloud	-0.222682
feels_like_celsius	-0.068775
feels_like_fahrenheit	-0.068819
visibility_km	-0.140602
visibility_miles	-0.153333
uv_index	-0.082052
gust_mph	-0.105203
gust_kph	-0.105229
air_quality_Carbon_Monoxide	0.471069
air_quality_Ozone	0.072649
air_quality_Nitrogen_dioxide	0.548457
air_quality_Sulphur_dioxide	0.283694
air_quality_PM2.5	0.769686
air_quality_PM10	0.513657
air_quality_us-epa-index	1.000000
air_quality_gb-defra-index	0.940809
moon_illumination	0.003262
Hour_sin	-0.013703
Hour_cos	-0.149723
Month_sin	0.159906
Month_cos	0.284928

	air_quality_gb-defra-index	
moon_illumination \		
latitude	0.061964	-
0.000289		
longitude	0.122535	-
0.001131		
temperature_celsius	-0.047209	-
0.014185		
temperature_fahrenheit	-0.047210	-
0.014187		
wind_mph	-0.058086	
0.008497		
wind_kph	-0.058119	
0.008491		
wind_degree	-0.000876	-
0.015165		
pressure_mb	0.048552	
0.004325		
pressure_in	0.048720	
0.004248		
precip_mm	-0.070285	
0.005482		

precip_in	-0.067839	
0.005246		
humidity	-0.196240	
0.007630		
cloud	-0.221123	
0.012135		
feels_like_celsius	-0.059620	-
0.011345		
feels_like_fahrenheit	-0.059669	-
0.011333		
visibility_km	-0.140553	-
0.002241		
visibility_miles	-0.154479	-
0.003847		
uv_index	-0.088459	-
0.018348		
gust_mph	-0.093380	
0.006626		
gust_kph	-0.093412	
0.006650		
air_quality_Carbon_Monoxide	0.433213	-
0.001637		
air_quality_Ozone	0.075244	-
0.003162		
air_quality_Nitrogen_dioxide	0.540302	
0.002133		
air_quality_Sulphur_dioxide	0.280532	-
0.001131		
air_quality_PM2.5	0.720709	-
0.002465		
air_quality_PM10	0.480097	-
0.001317		
air_quality_us-epa-index	0.940809	
0.003262		
air_quality_gb-defra-index	1.000000	
0.000157		
moon_illumination	0.000157	
1.000000		
Hour_sin	-0.015624	
0.006133		
Hour_cos	-0.135998	
0.006608		
Month_sin	0.163650	
0.153978		
Month_cos	0.272945	-
0.055846		

Hour_sin	Hour_cos	Month_sin	Month_cos
----------	----------	-----------	-----------

latitude	-0.117487	-0.291952	0.000908	-0.000320
longitude	-0.703979	0.140608	-0.001179	0.003569
temperature_celsius	-0.233213	0.076940	-0.229576	-0.354801
temperature_fahrenheit	-0.233204	0.076925	-0.229575	-0.354802
wind_mph	-0.063190	-0.029292	0.007970	-0.043862
wind_kph	-0.063147	-0.029262	0.007973	-0.043886
wind_degree	-0.171327	-0.113386	-0.038939	-0.066503
pressure_mb	0.110166	-0.116270	0.068977	0.124114
pressure_in	0.109265	-0.115996	0.068855	0.124354
precip_mm	-0.044366	0.076860	-0.018893	-0.028547
precip_in	-0.044819	0.074977	-0.018558	-0.028431
humidity	0.362730	0.371332	0.024546	0.132608
cloud	0.062810	0.165537	0.002372	0.050845
feels_like_celsius	-0.212557	0.128124	-0.224689	-0.353872
feels_like_fahrenheit	-0.212563	0.128146	-0.224698	-0.353886
visibility_km	-0.077997	0.142437	-0.042055	-0.066737
visibility_miles	-0.082728	0.136490	-0.048847	-0.071232
uv_index	-0.203808	-0.488760	-0.164533	-0.349198
gust_mph	-0.057325	0.060583	0.006230	-0.094471
gust_kph	-0.057360	0.060566	0.006216	-0.094514
air_quality_Carbon_Monoxide	-0.079739	0.061676	0.052961	0.056585
air_quality_Ozone	-0.258583	-0.174392	0.019732	-0.044031
air_quality_Nitrogen_dioxide	-0.050640	-0.048762	0.096101	0.190322
air_quality_Sulphur_dioxide	-0.065229	-0.014875	0.042896	0.090919
air_quality_PM2.5	-0.005076	-0.080553	0.104827	0.160459
air_quality_PM10	0.009843	-0.103605	0.072800	0.106816

air_quality_us-epa-index	-0.013703	-0.149723	0.159906	0.284928
air_quality_gb-defra-index	-0.015624	-0.135998	0.163650	0.272945
moon_illumination	0.006133	0.006608	0.153978	-0.055846
Hour_sin	1.000000	-0.030083	0.055253	0.308274
Hour_cos	-0.030083	1.000000	0.003468	-0.116881
Month_sin	0.055253	0.003468	1.000000	0.148197
Month_cos	0.308274	-0.116881	0.148197	1.000000

[33 rows x 33 columns]


#Plotting Matrix

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(correlation_matrix, annot=False, cmap='coolwarm',  
fmt=".2f", linewidths=0.5)
```

```
plt.title("Numerical Inputs Correlation Matrix")
```

```
plt.show()
```

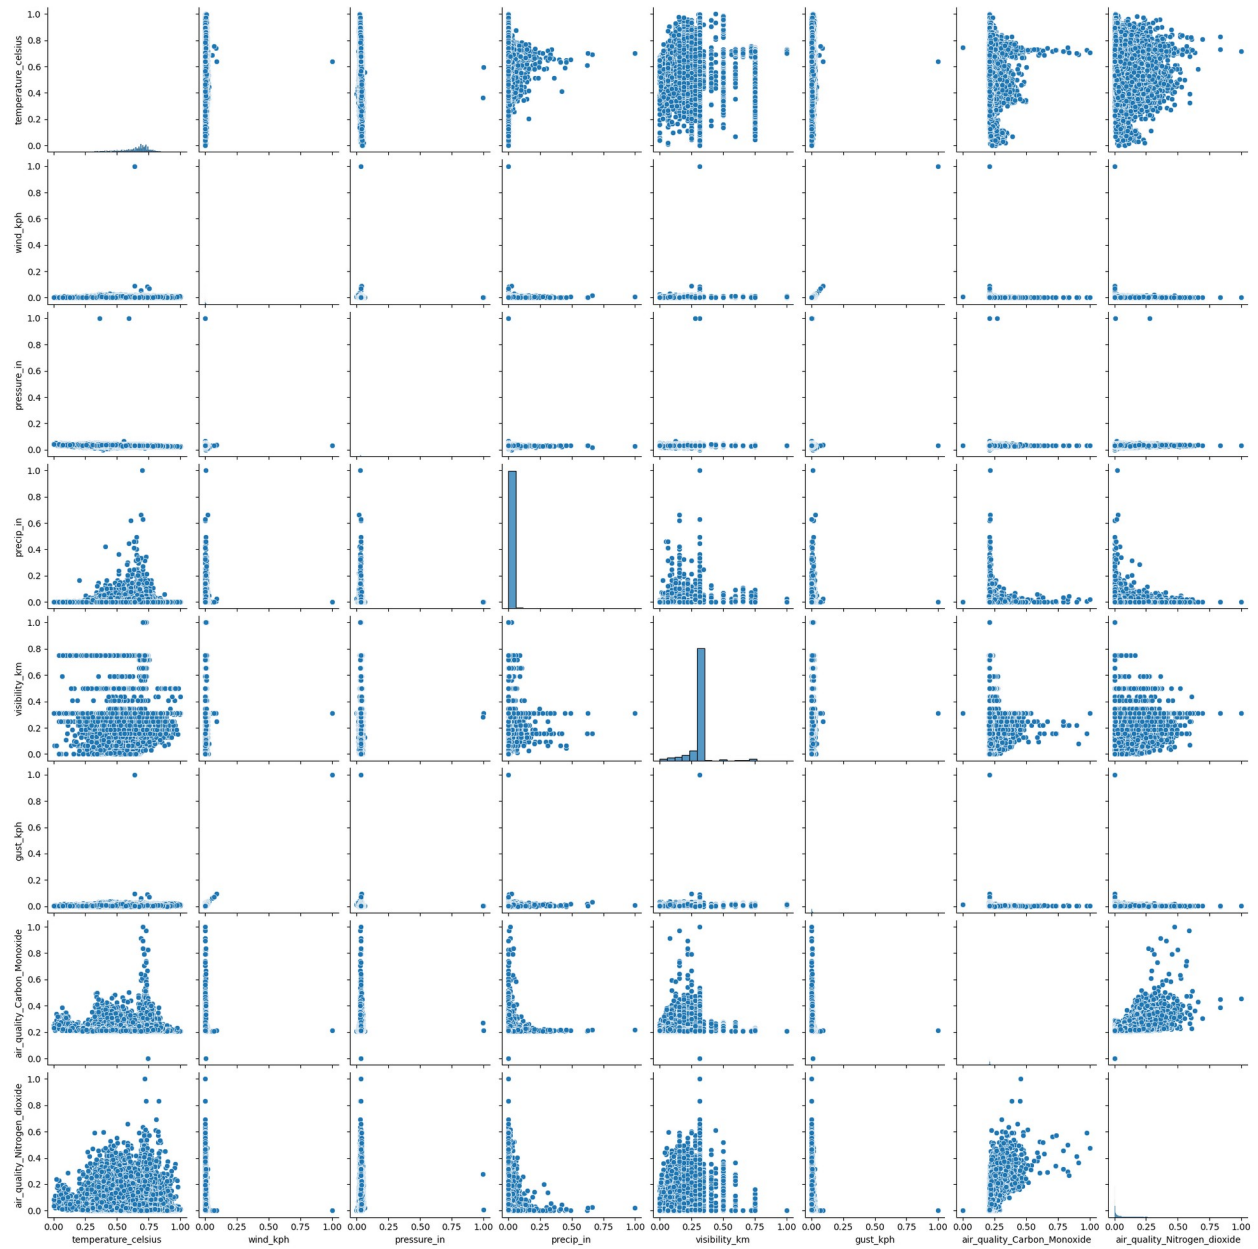



Aside from the variables in different measures, we can see that the air quality between some components has high correlations; still, I don't believe it is good to remove those.

#Plotting some Interest Variables to get relations

```
sns.pairplot(df_scaled[['temperature_celsius', 'wind_kph', 'pressure_in',
'precip_in', 'visibility_km', 'gust_kph', 'air_quality_Carbon_Monoxide',
'air_quality_Nitrogen_dioxide']])

plt.show()
```




some interest variables relations some of the with linear relation

#Plotting boxplots for temperature and condition

```
plt.figure(figsize=(12, 6))
sns.boxplot(
    x=df['condition_text'],
    y=df['temperature_celsius'],
    hue=df['condition_text'],
    palette="coolwarm",
    legend=False
)
```

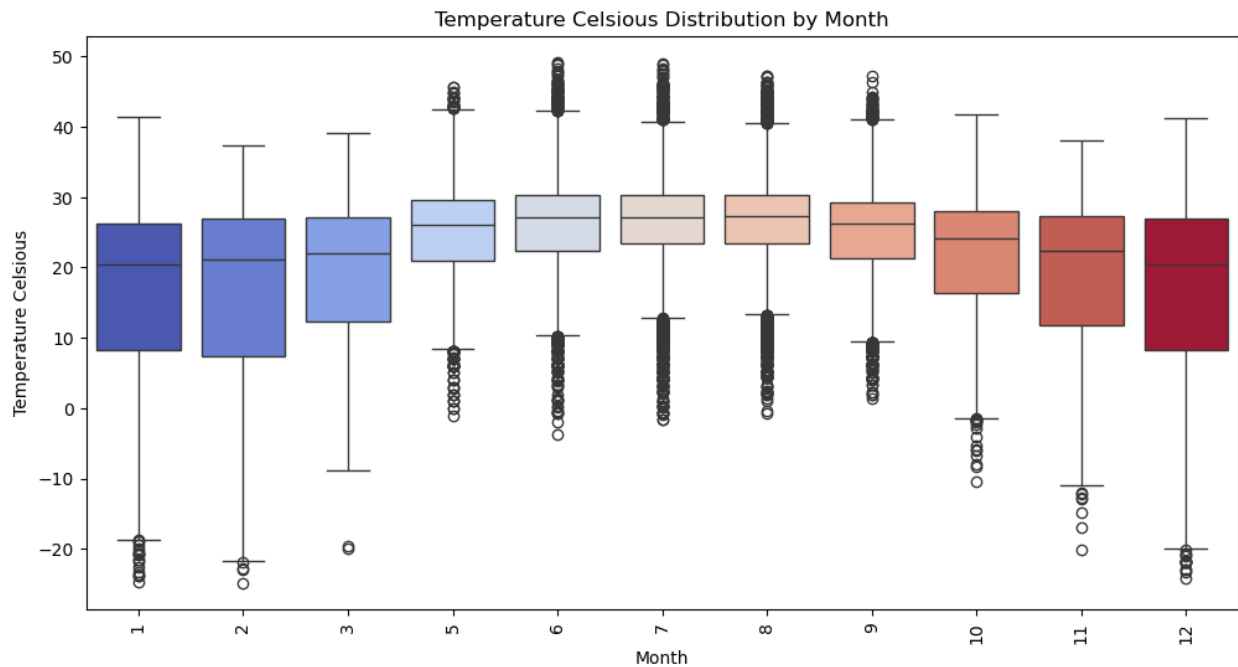
```
plt.xticks(rotation=90)
plt.xlabel("Condition")
plt.ylabel("Temperature Celsious")
plt.title("Temperature Celsious Distribution by Condition")
plt.show()
```



```
plt.figure(figsize=(12, 6))
sns.boxplot(
    x=df['month'],
    y=df['temperature_celsius'],
    hue=df['month'],
    palette="coolwarm",
    legend=False
)

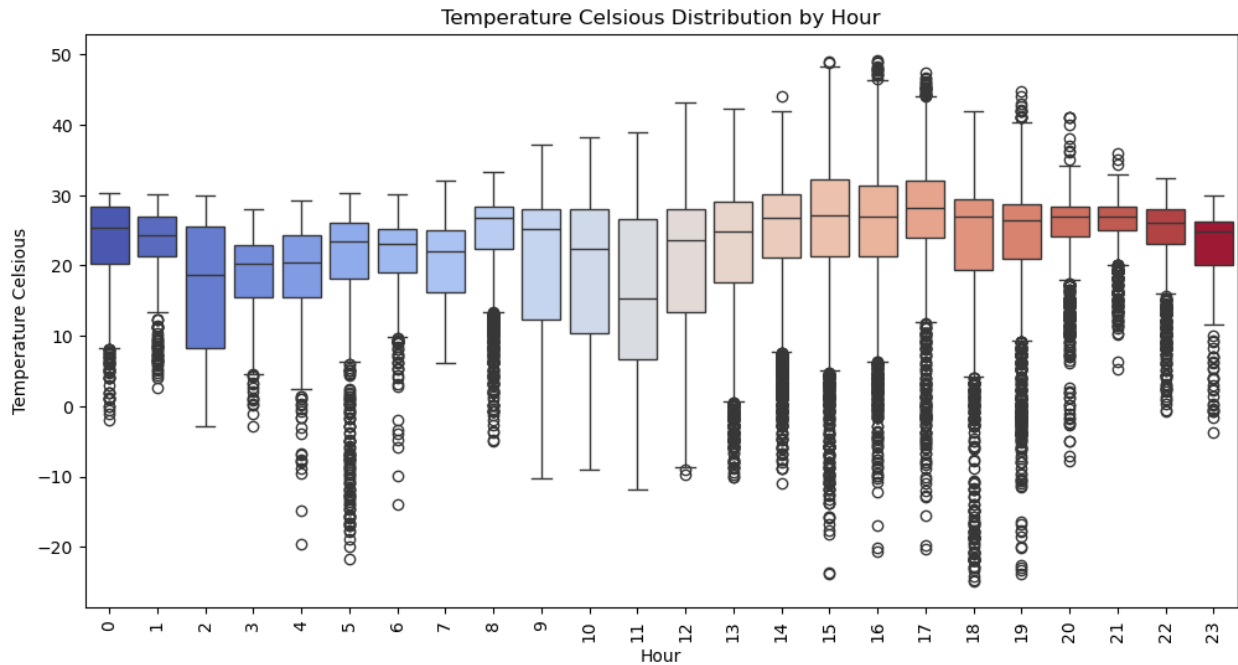
plt.xticks(rotation=90)
plt.xlabel("Month")
```

```
plt.ylabel("Temperature Celsius")
plt.title("Temperature Celsius Distribution by Month")
plt.show()
```



```
plt.figure(figsize=(12, 6))
sns.boxplot(
    x=df['hour'],
    y=df['temperature_celsius'],
    hue=df['hour'],
    palette="coolwarm",
    legend=False
)

plt.xticks(rotation=90)
plt.xlabel("Hour")
plt.ylabel("Temperature Celsius")
plt.title("Temperature Celsius Distribution by Hour")
plt.show()
```



#Plotting boxplots for temperature and condition

```
plt.figure(figsize=(12, 6))
```

```
sns.boxplot(
    x=df['condition_text'],
    y=df['precip_in'],
    hue=df['condition_text'],
    legend=False
)
```


```
plt.xticks(rotation=90)
```

```
plt.xlabel("Condition")
```

```
plt.ylabel("Precipitation Inches")
```

```
plt.title("Precipitation Inches Distribution by Condition")
```

```
plt.show()
```

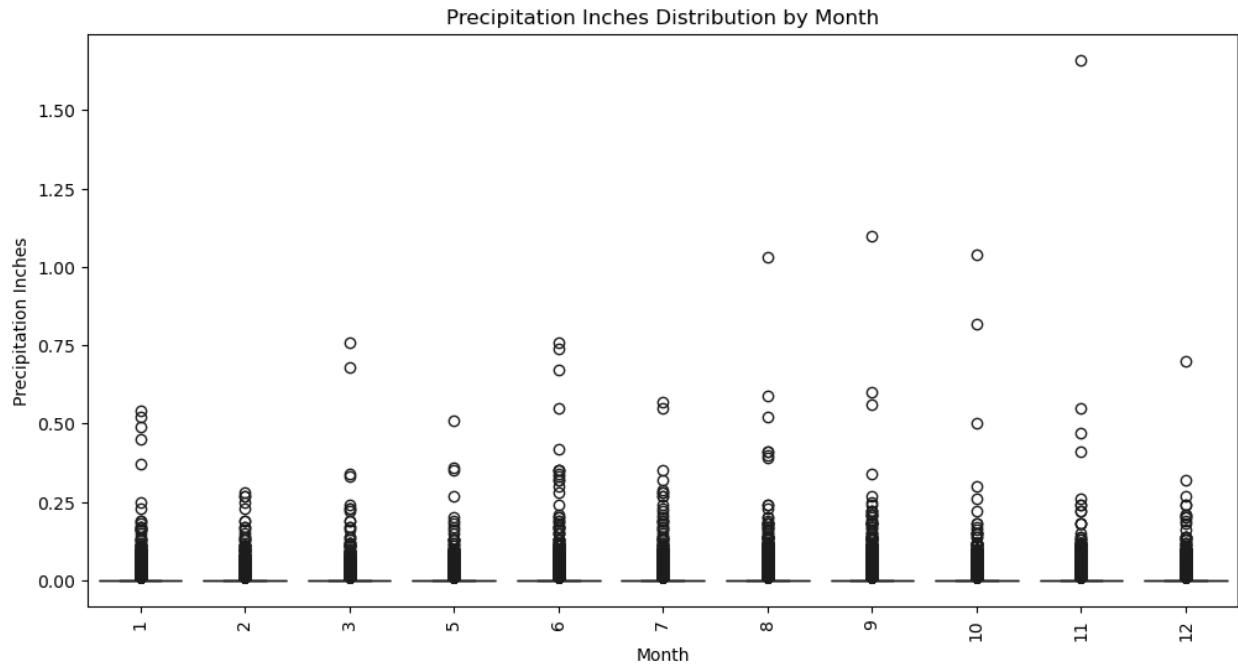


#Plotting boxplots for temperature and condition

```
plt.figure(figsize=(12, 6))
```

```
sns.boxplot(
    x=df['month'],
    y=df['precip_in'],
    hue=df['month'],
    legend=False
)
```

```
plt.xticks(rotation=90)
plt.xlabel("Month")
plt.ylabel("Precipitation Inches")
plt.title("Precipitation Inches Distribution by Month")
plt.show()
```



#Plotting boxplots for temperature and condition

```
plt.figure(figsize=(12, 6))
```

```
sns.boxplot(  
    x=df['hour'],  
    y=df['precip_in'],  
    hue=df['hour'],  
    legend=False  
)
```

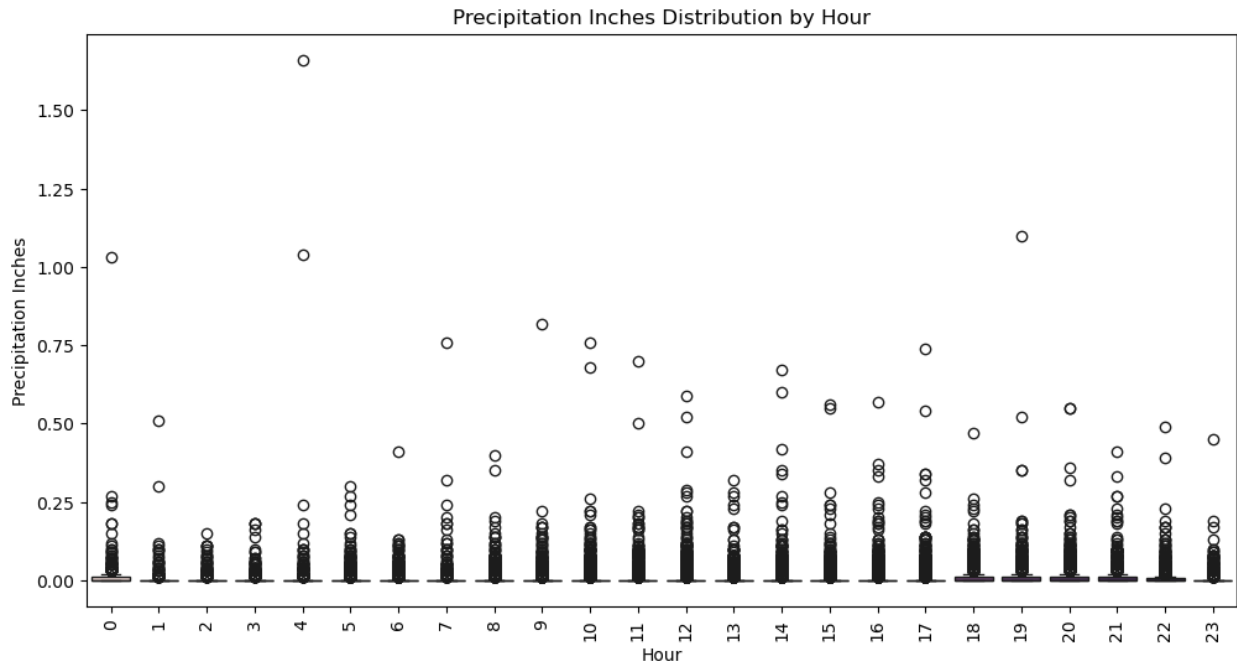
```
plt.xticks(rotation=90)
```

```
plt.xlabel("Hour")
```

```
plt.ylabel("Precipitation Inches")
```

```
plt.title("Precipitation Inches Distribution by Hour")
```

```
plt.show()
```



For the modeling part, I only did it thinking of one particular location because the way the dataset is formatted in one day can have multiple measures for distinct locations, and we know that the location, of course, affects the weather. In this case I chose Tegucigalpa, Honduras.

#Setting a location for prediction analysis

```
location='Tegucigalpa'
df_scaled_location= df_scaled[df_scaled['location_name']==location]
df_scaled_location= df_scaled_location.sort_values(by='last_updated')
df_scaled_location
```

	country	location_name	timezone	last_updated
74	Honduras	Tegucigalpa	America/Tegucigalpa	2024-05-16 02:45:00
269	Honduras	Tegucigalpa	America/Tegucigalpa	2024-05-16 08:00:00
464	Honduras	Tegucigalpa	America/Tegucigalpa	2024-05-17 10:00:00
656	Honduras	Tegucigalpa	America/Tegucigalpa	2024-05-18 08:30:00
850	Honduras	Tegucigalpa	America/Tegucigalpa	2024-05-19 08:15:00
...
58732	Honduras	Tegucigalpa	America/Tegucigalpa	2025-03-15 03:45:00
58927	Honduras	Tegucigalpa	America/Tegucigalpa	2025-03-16 03:45:00

59122	Honduras	Tegucigalpa	America/Tegucigalpa	2025-03-17 03:45:00
59317	Honduras	Tegucigalpa	America/Tegucigalpa	2025-03-18 03:45:00
59512	Honduras	Tegucigalpa	America/Tegucigalpa	2025-03-19 03:30:00

	condition_text	wind_direction	sunrise	sunset	moonrise
moonset ... \					
74	Partly cloudy	WSW	5	18	12.0
0.0 ...					
269	Partly cloudy	NNE	5	18	12.0
0.0 ...					
464	Partly cloudy	NNW	5	18	13.0
1.0 ...					
656	Partly cloudy	N	5	18	14.0
2.0 ...					
850	Overcast	N	5	18	15.0
2.0 ...					
...
...					
58732	Partly cloudy	SSW	5	17	19.0
6.0 ...					
58927	Clear	SSW	5	17	20.0
7.0 ...					
59122	Partly cloudy	N	5	18	20.0
7.0 ...					
59317	Partly cloudy	N	5	18	21.0
8.0 ...					
59512	Partly cloudy	NNE	5	18	22.0
9.0 ...					

	air_quality_Sulphur_dioxide	air_quality_PM2.5	air_quality_PM10
\			
74	0.950579	0.011658	0.237583
269	0.950541	0.011658	0.238103
464	0.950617	0.001372	0.234844
656	0.950788	0.004223	0.235884
850	0.950569	0.010481	0.237469
...
58732	0.951993	0.013641	0.237284
58927	0.952239	0.015131	0.237636

59122	0.951360	0.018914	0.238387
59317	0.950674	0.006648	0.235923
59512	0.950709	0.006075	0.235806

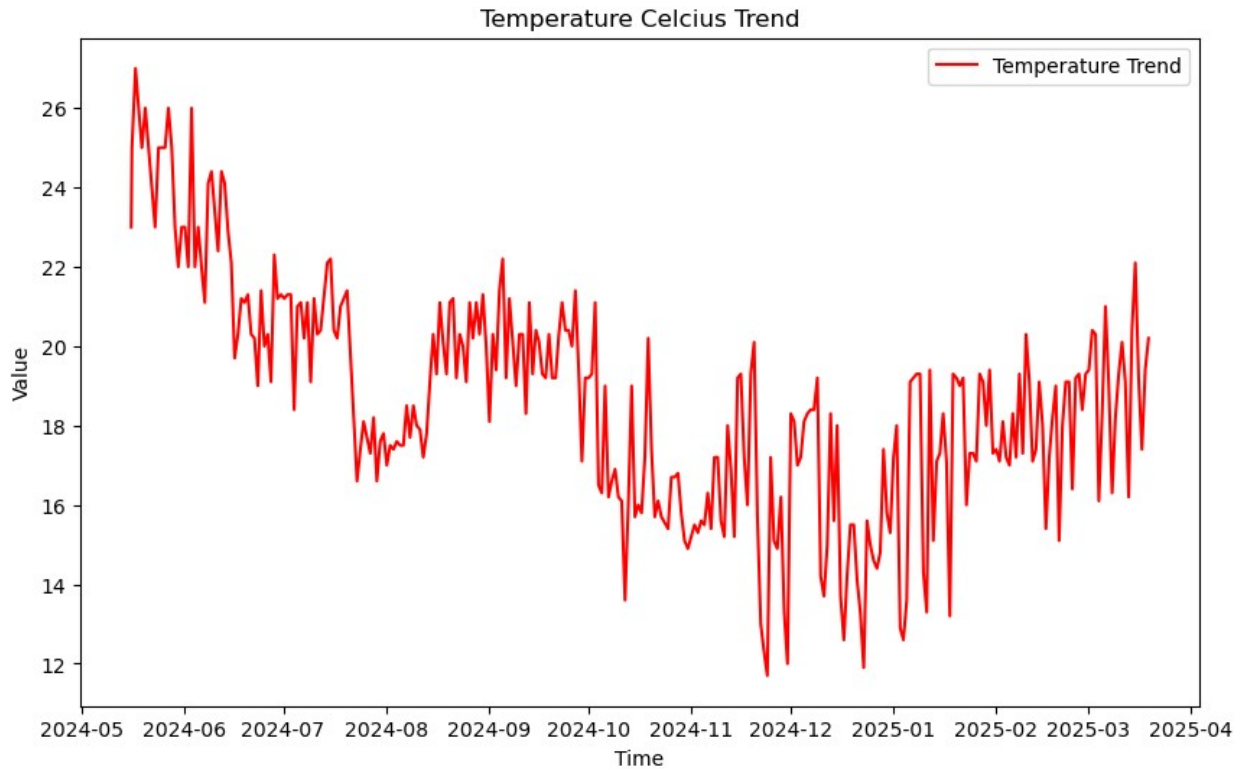
	air_quality_us-epa-index	air_quality_gb-defra-index	\
74	0.2	0.111111	
269	0.2	0.111111	
464	0.0	0.000000	
656	0.0	0.000000	
850	0.2	0.111111	
...
58732	0.2	0.111111	
58927	0.2	0.222222	
59122	0.2	0.222222	
59317	0.0	0.000000	
59512	0.0	0.000000	

	moon_illumination	Hour_sin	Hour_cos	Month_sin	Month_cos
74	0.55	0.750000	0.933013	0.75	0.066987
269	0.55	0.933013	0.250000	0.75	0.066987
464	0.64	0.750000	0.066987	0.75	0.066987
656	0.73	0.933013	0.250000	0.75	0.066987
850	0.81	0.933013	0.250000	0.75	0.066987
...
58732	0.99	0.853553	0.853553	1.00	0.500000
58927	0.96	0.853553	0.853553	1.00	0.500000
59122	0.92	0.853553	0.853553	1.00	0.500000
59317	0.86	0.853553	0.853553	1.00	0.500000
59512	0.79	0.853553	0.853553	1.00	0.500000

[303 rows x 45 columns]

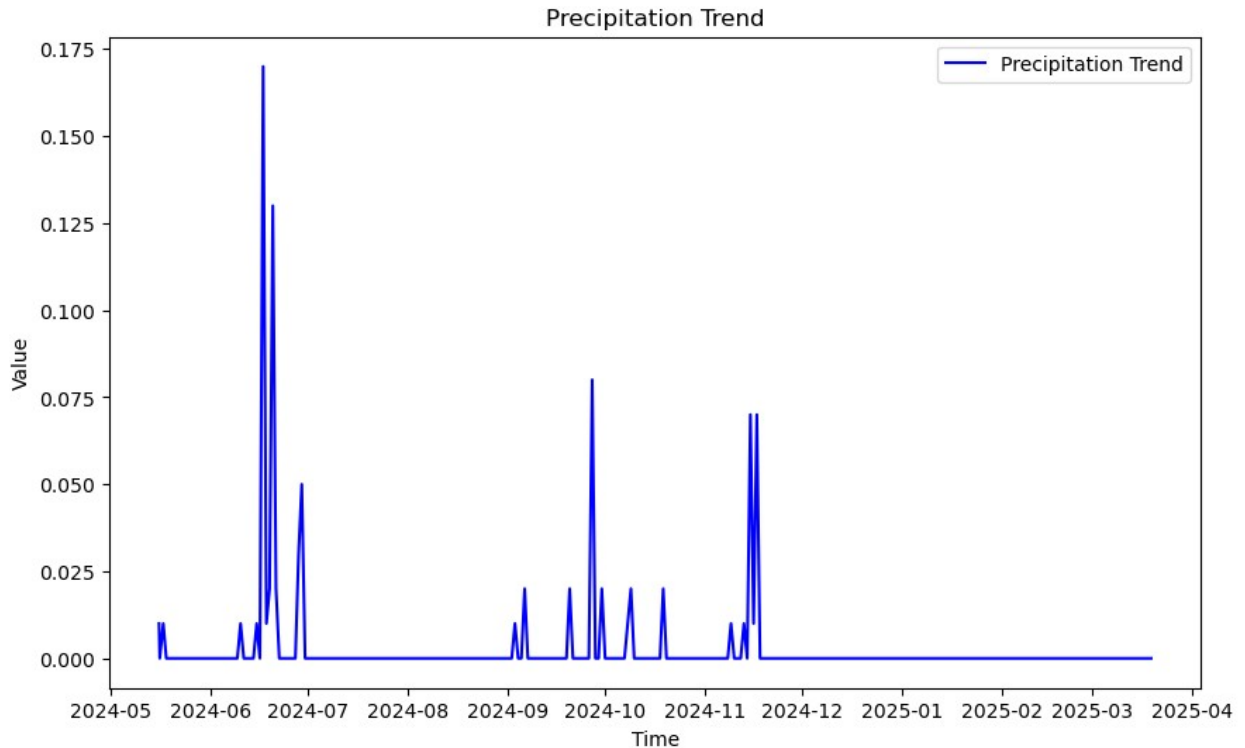
#Plotting Temperature trend for Location

```
plt.figure(figsize=(10,6))
plt.plot(df_scaled_location['last_updated'],scaler.inverse_transform(d
f_scaled_location[columns])[:,columns.index("temperature_celsius")],
label="Temperature Trend", color="red")
plt.xlabel("Time")
plt.ylabel("Value")
plt.title("Temperature Celcius Trend")
plt.legend()
plt.show()
```



#Plotting Precipitation trend for Location

```
plt.figure(figsize=(10,6))
plt.plot(df_scaled_location['last_updated'],scaler.inverse_transform(d
f_scaled_location[columns])[:,columns.index("precip_in")],
label="Precipitation Trend", color="blue")
plt.xlabel("Time")
plt.ylabel("Value")
plt.title("Precipitation Trend")
plt.legend()
plt.show()
```



The first model that I will use is an LSTM neural network because I am used to working with it and looking at the plots for precipitation and temperature, a complex model is needed

Method to create sequences for LSTM Model

```
def create_sequences(data, time_steps=5):
    X, y = [], []
    for i in range(len(data) - time_steps):
        X.append(data[i:i+time_steps])
        y.append(data[i+time_steps])
    return np.array(X), np.array(y)
```

#Creating sequences for LSTM

```
X, y = create_sequences(df_scaled_location[columns].to_numpy())
```

#Splitting data in training and validation

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.2, shuffle=False)
```

#creating model structure

```
model = Sequential()
model.add(LSTM(units=248, return_sequences=False,
    input_shape=(X_train.shape[1], X_train.shape[2])))
model.add(Dense(units=128))
model.add(Dense(units=64))
model.add(Dense(units=len(columns)))
```

```

model.compile(optimizer='adam', loss='mean_squared_error')

#Repeating the dataset to improve training
train_ds = tf.data.Dataset.from_tensor_slices(
    (X_train, y_train)).shuffle(len(X_train)).repeat().batch(32)
test_ds = tf.data.Dataset.from_tensor_slices((X_test,
y_test)).batch(32)

# Train the model
history = model.fit(train_ds, epochs=20, batch_size=32,
validation_data=(X_test, y_test), steps_per_epoch=500)

Epoch 1/20
500/500 [=====] - 10s 15ms/step - loss:
0.0099 - val_loss: 0.0083
Epoch 2/20
500/500 [=====] - 6s 12ms/step - loss: 0.0066
- val_loss: 0.0096
Epoch 3/20
500/500 [=====] - 7s 14ms/step - loss: 0.0054
- val_loss: 0.0126
Epoch 4/20
500/500 [=====] - 6s 13ms/step - loss: 0.0036
- val_loss: 0.0161
Epoch 5/20
500/500 [=====] - 4s 7ms/step - loss: 0.0020
- val_loss: 0.0253
Epoch 6/20
500/500 [=====] - 7s 14ms/step - loss: 0.0011
- val_loss: 0.0323
Epoch 7/20
500/500 [=====] - 6s 13ms/step - loss:
5.8780e-04 - val_loss: 0.0368
Epoch 8/20
500/500 [=====] - 5s 10ms/step - loss:
3.8987e-04 - val_loss: 0.0308
Epoch 9/20
500/500 [=====] - 3s 7ms/step - loss:
2.9084e-04 - val_loss: 0.0327
Epoch 10/20
500/500 [=====] - 3s 7ms/step - loss:
2.5168e-04 - val_loss: 0.0301
Epoch 11/20
500/500 [=====] - 3s 7ms/step - loss:
1.7950e-04 - val_loss: 0.0312
Epoch 12/20
500/500 [=====] - 3s 7ms/step - loss:
1.7881e-04 - val_loss: 0.0279
Epoch 13/20
500/500 [=====] - 3s 7ms/step - loss:

```

```

1.5592e-04 - val_loss: 0.0268
Epoch 14/20
500/500 [=====] - 3s 7ms/step - loss:
1.7166e-04 - val_loss: 0.0269
Epoch 15/20
500/500 [=====] - 4s 7ms/step - loss:
1.0265e-04 - val_loss: 0.0246
Epoch 16/20
500/500 [=====] - 3s 7ms/step - loss:
1.2973e-04 - val_loss: 0.0252
Epoch 17/20
500/500 [=====] - 3s 7ms/step - loss:
8.7254e-05 - val_loss: 0.0236
Epoch 18/20
500/500 [=====] - 3s 7ms/step - loss:
1.0811e-04 - val_loss: 0.0238
Epoch 19/20
500/500 [=====] - 3s 7ms/step - loss:
9.0295e-05 - val_loss: 0.0222
Epoch 20/20
500/500 [=====] - 4s 7ms/step - loss:
8.2237e-05 - val_loss: 0.0201

```

Training goes smoothly with good loss measures

#Predicting Fit Values

```
train_predictions = model.predict(X_train)
```

Reverse scaling to get back original values

```

train_predictions_rescaled =
scaler.inverse_transform(train_predictions)
y_train_rescaled = scaler.inverse_transform(y_train)

```

Calculate metrics

```

mse_train = mean_squared_error(y_train_rescaled,
train_predictions_rescaled)
rmse_train = np.sqrt(mse_train)
mae_train = mean_absolute_error(y_train_rescaled,
train_predictions_rescaled)

```

Print results

```

print(f"Train Mean Squared Error (MSE): {mse_train:.4f}")
print(f"Train Root Mean Squared Error (RMSE): {rmse_train:.4f}")
print(f"Train Mean Absolute Error (MAE): {mae_train:.4f}")

```


```

8/8 [=====] - 0s 3ms/step
Train Mean Squared Error (MSE): 1362.2510
Train Root Mean Squared Error (RMSE): 36.9087
Train Mean Absolute Error (MAE): 9.2302


```

In this case we used 3 error measures mse rmse and mae

```
#Plot Fit vs Actual for temperature
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(y_train_rescaled):],
y_train_rescaled[:,columns.index("temperature_celsius")],
label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(train_predictions_rescaled):], train_predictions_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("LSTM Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



```
#Plot Fit vs Actual for temperature
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(y_train_rescaled):],
y_train_rescaled[:,columns.index("precip_in")], label="Actual
Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(train_predictions_rescaled):], train_predictions_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("LSTM Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



For what we can see in the plot the model does a good job fitting, would require more training for it to fit perfectly but that could not be good due to overfitting.

#Predicting Test Values

```
test_predictions = model.predict(X_test)
```

Reverse scaling to get back original values

```
test_predictions_rescaled = scaler.inverse_transform(test_predictions)
y_test_rescaled = scaler.inverse_transform(y_test)
```

Calculate metrics

```
mse_test = mean_squared_error(y_test_rescaled,
test_predictions_rescaled)
rmse_test = np.sqrt(mse_train)
mae_test = mean_absolute_error(y_test_rescaled,
test_predictions_rescaled)
r2_test = r2_score(y_test_rescaled, test_predictions_rescaled)
```

Print results


```
print(f"Test Mean Squared Error (MSE): {mse_test:.4f}")
print(f"Test Root Mean Squared Error (RMSE): {rmse_test:.4f}")
print(f"Test Mean Absolute Error (MAE): {mae_test:.4f}")
```

```
2/2 [=====] - 0s 11ms/step
Test Mean Squared Error (MSE): 8681.2693
Test Root Mean Squared Error (RMSE): 36.9087
Test Mean Absolute Error (MAE): 24.5857
```

Considerable Higher measures for the validation set, not as much to call it overfitting.



```
#Plot Predicted vs Actual for temperature
```

```
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(y_test_rescaled):],
y_test_rescaled[:, columns.index("temperature_celsius")],
label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(test_predictions_rescaled):], test_predictions_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("LSTM Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



```
#Plot Predicted vs Actual for temperature
```

```
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(y_test_rescaled):],
y_test_rescaled[:, columns.index("precip_in")], label="Actual
Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(test_predictions_rescaled):], test_predictions_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("LSTM Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



We can see that the model tries to simulate the complex behaviour of temperature but struggles to do so. other reason could be that Im trying to predict multiple measures at the same time, maybe just focusing on temperature or precipitation alone would deliver better results.

The next model to implement would be Random Forest

Same process of split, train transform to the original scale, calculating error measures and plotting

```
#Preparing data for Random Forest Model
X_flat = X.reshape((X.shape[0], X.shape[1] * X.shape[2]))

#Training and test Split
X_train, X_test, y_train, y_test = train_test_split(X_flat, y,
test_size=0.2, shuffle=False)

#Model Fit
rf_model = RandomForestRegressor(n_estimators=100)
rf_model.fit(X_train, y_train)

#Prediction Fit
rf_train_preds = rf_model.predict(X_train)

# Reverse scaling to get back original values
rf_train_preds_rescaled = scaler.inverse_transform(rf_train_preds)
rfy_train_rescaled = scaler.inverse_transform(y_train)

#Prediction Test
rf_test_preds = rf_model.predict(X_test)

# Reverse scaling to get back original values
```

```
rf_test_preds_rescaled = scaler.inverse_transform(rf_test_preds)
rfy_test_rescaled = scaler.inverse_transform(y_test)
```

Error Measures

```
mse_train_rf = mean_squared_error(rfy_train_rescaled,
rf_train_preds_rescaled)
rmse_train_rf = np.sqrt(mse_train_rf)
train_mae_rf = mean_absolute_error(rfy_train_rescaled,
rf_train_preds_rescaled)
mse_test_rf = mean_squared_error(rfy_test_rescaled,
rf_test_preds_rescaled)
rmse_test_rf = np.sqrt(mse_test_rf)
test_mae_rf = mean_absolute_error(rfy_test_rescaled,
rf_test_preds_rescaled)


print(f"Random Forest Train MSE: {mse_train_rf:.4f}")
print(f"Random Forest Train RMSE: {rmse_train_rf:.4f}")
print(f"Random Forest Train MAE: {train_mae_rf:.4f}")
print(f"Random Forest Test MSE: {mse_test_rf:.4f}")
print(f"Random Forest Test RMSE: {rmse_test_rf:.4f}")
print(f"Random Forest Test MAE: {test_mae_rf:.4f}")
```

```
Random Forest Train MSE: 178.4549
Random Forest Train RMSE: 13.3587
Random Forest Train MAE: 3.1959
Random Forest Test MSE: 2548.6511
Random Forest Test RMSE: 50.4842
Random Forest Test MAE: 7.8328
```


To my surprise the random forest model has better error measures

Plotting Prediction of fit values vs actual values

```
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("temperature_celsius")], label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(rf_train_preds_rescaled):], rf_train_preds_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Random Forest Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```




```
# Plotting Prediction of fit values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("precip_in")], label="Actual Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(rf_train_preds_rescaled):], rf_train_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Random Forest Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```




It does a really good job fitting the model

```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("temperature_celsius")],
label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(rf_test_preds_rescaled):], rf_test_preds_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Random Forest Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("precip_in")], label="Actual
Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(rf_test_preds_rescaled):], rf_test_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Random Forest Forecast: Predictions vs Actual Values")
plt.legend()
plt.show()
```



For this model it is closer the prediction to the actual value, but it doesn't try to simulate the complex trend of the measures.

No I will combine the 2 models using a stack model in this case a linear regression to see if I can improve the performance.

Training an Ensemble of models stacking them

```
X_train_stack = np.column_stack((train_predictions, rf_train_preds))
X_test_stack = np.column_stack((test_predictions, rf_test_preds))
```

#Using Linear Regresion as Meta Model to combine predictions

```
meta_model = LinearRegression()
meta_model.fit(X_train_stack, y_train)
```

Predict using the stacked model

```
stacked_train_preds = meta_model.predict(X_train_stack)
stacked_test_preds = meta_model.predict(X_test_stack)
```

Reverse scaling to get back original values

```
stacked_train_preds_rescaled =
scaler.inverse_transform(stacked_train_preds)
stacked_test_preds_rescaled =
scaler.inverse_transform(stacked_test_preds)
```

Evaluate the performance of the stacked model

```
train_mse = mean_squared_error(rfy_train_rescaled,
stacked_train_preds_rescaled)
train_rmse = np.sqrt(mse_train_rf)
train_mae = mean_absolute_error(rfy_train_rescaled,
stacked_train_preds_rescaled)
```

```

test_mse = mean_squared_error(rfy_test_rescaled,
                              stacked_test_preds_rescaled)
test_rmse = np.sqrt(test_mse)
test_mae = mean_absolute_error(rfy_test_rescaled,
                              stacked_test_preds_rescaled)

print(f"Stacked Model Train MSE: {train_mse:.4f}")
print(f"Stacked Model Train RMSE: {train_rmse:.4f}")
print(f"Stacked Model Train MAE: {train_mae:.4f}")
print(f"Stacked Model Test MSE: {test_mse:.4f}")
print(f"Stacked Model Test RMSE: {test_rmse:.4f}")
print(f"Stacked Model Test MAE: {test_mae:.4f}")

Stacked Model Train MSE: 11.9690
Stacked Model Train RMSE: 13.3587
Stacked Model Train MAE: 0.8627
Stacked Model Test MSE: 5177.3622
Stacked Model Test RMSE: 71.9539
Stacked Model Test MAE: 16.2126


```

Error measures in a good area mainly on the training set on the validation still the random forest seem to perform better.


```

# Plotting Prediction of fit values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("temperature_celsius")], label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_train_preds_rescaled):], stacked_train_preds_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Stacked Model(Linear Regression): Temperature Fit
Predictions vs Actual Values")
plt.legend()
plt.show()

```





```
# Plotting Prediction of fit values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("precip_in")], label="Actual Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_train_preds_rescaled):], stacked_train_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Stacked Model(Linear Regression): Precipitation Fit
Predictions vs Actual Values")
plt.legend()
plt.show()
```




We can see that the model almost perfectly fits the training data

```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("temperature_celsius")],
label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_test_preds_rescaled):], stacked_test_preds_rescaled[:, 2],
label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Stacked Model(Linear Regression): Temperature Test
Predictions vs Actual Values")
plt.legend()
plt.show()
```



```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("precip_in")], label="Actual
Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_test_preds_rescaled):], stacked_test_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Stacked Model(Linear Regression): Precipitation Test
Predictions vs Actual Values")
plt.legend()
plt.show()
```



The model in comparison with the LSTM alone it got closer to the prediction, the random forest still performs better but this one tries to simulate those high and low points also.

Now instead of using a linear regression as the stacked model, I used the Random Forest again.

```
# Training an Ensemble of models stacking them
```

```
X_train_stack = np.column_stack((train_predictions, rf_train_preds))
X_test_stack = np.column_stack((test_predictions, rf_test_preds))
```

```
#Using RandomForest as Meta Model to combine predictions
```

```
meta_model = RandomForestRegressor(n_estimators=50)
meta_model.fit(X_train_stack, y_train)
```

```
# Predict using the stacked model
```

```
stacked_train_preds = meta_model.predict(X_train_stack)
stacked_test_preds = meta_model.predict(X_test_stack)
```

```
# Reverse scaling to get back original values
```

```
stacked_train_preds_rescaled =
scaler.inverse_transform(stacked_train_preds)
stacked_test_preds_rescaled =
scaler.inverse_transform(stacked_test_preds)
```

```
# Evaluate the performance of the stacked model
```

```
train_mse = mean_squared_error(rfy_train_rescaled,
stacked_train_preds_rescaled)
train_rmse = np.sqrt(mse_train_rf)
train_mae = mean_absolute_error(rfy_train_rescaled,
```

```

stacked_train_preds_rescaled)
test_mse = mean_squared_error(rfy_test_rescaled,
stacked_test_preds_rescaled)
test_rmse = np.sqrt(test_mse)
test_mae = mean_absolute_error(rfy_test_rescaled,
stacked_test_preds_rescaled)

print(f"Stacked Model Train MSE: {train_mse:.4f}")
print(f"Stacked Model Train RMSE: {train_rmse:.4f}")
print(f"Stacked Model Train MAE: {train_mae:.4f}")
print(f"Stacked Model Test MSE: {test_mse:.4f}")
print(f"Stacked Model Test RMSE: {test_rmse:.4f}")
print(f"Stacked Model Test MAE: {test_mae:.4f}")

Stacked Model Train MSE: 94.3645
Stacked Model Train RMSE: 13.3587
Stacked Model Train MAE: 1.9406
Stacked Model Test MSE: 2696.4991
Stacked Model Test RMSE: 51.9278
Stacked Model Test MAE: 7.7604


```

Good Error Measures slightly higher in the validation set than the first random forest


```

# Plotting Prediction of fit values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("temperature_celsius")], label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_train_preds_rescaled):], stacked_train_preds_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Stacked Model(Random Forest): Temperature Fit Predictions
vs Actual Values")
plt.legend()
plt.show()


```




```
# Plotting Prediction of fit values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-
len(rfy_train_rescaled):], rfy_train_rescaled[:,
columns.index("precip_in")], label="Actual Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_train_preds_rescaled):], stacked_train_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Stacked Model(Random Forest): Precipitation Fit Predictions
vs Actual Values")
plt.legend()
plt.show()
```



```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("temperature_celsius")],
label="Actual Temperature")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_test_preds_rescaled):], stacked_test_preds_rescaled[:,
columns.index("temperature_celsius")], label="Predicted Temperature")
plt.xlabel("Dates")
plt.ylabel("Temperature")
plt.title("Stacked Model(Random Forest): Temperature Test Predictions
vs Actual Values")
plt.legend()
plt.show()
```



```
# Plotting Prediction of test values vs actual values
plt.figure(figsize=(12,6))
plt.plot(df_scaled_location['last_updated'][-len(rfy_test_rescaled):],
rfy_test_rescaled[:, columns.index("precip_in")], label="Actual
Precipitation")
plt.plot(df_scaled_location['last_updated'][-
len(stacked_test_preds_rescaled):], stacked_test_preds_rescaled[:,
columns.index("precip_in")], label="Predicted Precipitation")
plt.xlabel("Dates")
plt.ylabel("Precipitation")
plt.title("Stacked Model(Random Forest): Precipitation Test
Predictions vs Actual Values")
plt.legend()
plt.show()
```

The model in my opinion is trying to catch some of the variation of the measures, i believe is the influence of the lstm in the model, next steps to improve the model would be to add more stacked models to the equation and try to tune hyperparameters for better performance.

Now I will show based on the final random forest model some variable importance for the model

```
# In this case cause the features were duplicated cause we have the
# predicted values for the lstm model and the random regression model
# Creating new_columns with both models features
new_columns = columns + [column + "_rf" for column in columns]

#Variable importance

feature_importance = meta_model.feature_importances_
feature_names = new_columns

# Create a DataFrame
importance_df = pd.DataFrame({"Feature": feature_names, "Importance":
feature_importance})
importance_df = importance_df.sort_values(by="Importance",
ascending=False)


# Display importance
importance_df.style

<pandas.io.formats.style.Styler at 0x220edc266a0>
```

We can see that moon illumination, wind measures also the month are the most important variables, which in logical terms make sense.

#Variable importance using SHAP

```
explainer = shap.Explainer(meta_model, X_train_stack)
shap_values = explainer(X_test_stack, check_additivity=False)
shap.summary_plot(shap_values, X_test_stack, new_columns,
class_names=columns, max_display=X_test_stack.shape[1])
```



Same as before moon illumination, wind measures also the month are the most important variables, which in logical terms make sense.